BrandGuard

Role-by-role guide to reviewing and approving brand content

What the app does
BrandGuard checks marketing copy and uploaded documents against approved brand guidance, facts, terminology, and accessibility rules. It shows evidence beside each finding and records human decisions through review and approval.

The approval path

Draft Creator prepares content	**Review** Manager routes work	**Sign-off** Reviewer checks risk	**Approved** Manager decides

Before you start

Run **python scripts/setup_local.py**, then **python scripts/run_local.py**, and open **http://localhost:8000**. Local setup enables the fictional demo workspace. Choose Admin, Manager, Reviewer, or Creator from the **Demo accounts** panel on the login page; the guide intentionally does not print demo passwords.

The app runs offline with deterministic mock review by default. Image/PDF extraction and the retrieval model warm in local background workers. A human must verify findings and make approval decisions.

Use the four role sections on the following pages as a five-minute click path. The final page explains the three Round 3 improvements and why they matter.

Creator — submit, fix, and resubmit

Creators own their drafts and finding decisions. They can request review after checking the current version; they cannot approve their own content.

1. **Create a check** — Choose **New Content Check**. Paste copy or upload TXT, DOCX, PDF, PNG, JPEG, or SVG. You can use the fictional samples in **samples/**. **Why it matters:** One entry point supports both fast copy review and format-aware document checks.
2. **Read each finding** — Open the cited source and compare its exact quote with the highlighted claim. PDF page highlights and image OCR boxes link findings back to the source. **Why it matters:** Evidence makes each recommendation auditable and easier to validate.
3. **Record your decision** — Choose **Accept Fix**, **Edit**, or **Dispute** with a reason. A saved fix creates a new version; run the check again before routing it. **Why it matters:** Human feedback improves future rewrite ranking while preserving the original audit trail.
4. **Request review** — When the current version is checked, choose **Request Manager Review**. Follow the status stepper as the item moves through review, sign-off, and approval. **Why it matters:** The manager and reviewer see the version you actually checked.
5. **Resubmit changes** — If changes are requested or the item is rejected, update the revision form, save, and check the new version before resubmitting. **Why it matters:** Every decision stays attached to a specific revision.

What you can track

My Work shows your submissions, current status, due-date context where assigned, notifications, and feedback metrics. **Certificates** contains approval records and their verification status.

Manager — route work and decide

Managers coordinate review, assign compliance specialists, and make the final approval decision after required sign-off.

1. **Open Command Center** — Select an item and compare the submission, open findings, and creator decisions. **Why it matters:** Keeps the approval decision grounded in the checked revision.
2. **Assign a reviewer** — For high-severity, escalated, or expired-fact findings, assign an active Reviewer and set a due date. **Why it matters:** A four-eyes gate requires compliance review before final approval.
3. **Inspect guideline health** — Open **Guideline Health**, sort by dispute friction, inspect repeated reasons, and use **Recompute now** for fresh data. **Why it matters:** Recurring disputes can reveal clauses that need clearer wording.
4. **Make the workflow decision** — Approve standard-risk content, request changes, or reject. High-risk work becomes approvable by the Manager only after Reviewer sign-off. **Why it matters:** Separates routine routing from specialist compliance sign-off.
5. **Export decisions** — Download a CSV or PDF approval report when you need a portable review record. **Why it matters:** Gives teams a shareable record of decisions and due dates.

Reviewer — verify escalations

Reviewers are compliance specialists. They do not replace the Manager and do not make the final approval decision.

1. **Open Compliance Desk** — Choose an assigned submission and inspect the claim, cited source excerpt, and retrieved evidence side by side. **Why it matters:** Checks whether the evidence actually supports the claim.
2. **Resolve escalations** — Confirm or dismiss an escalated finding with a written justification. Add evidence notes when needed. **Why it matters:** Creates an explanation that the Creator and Manager can review.
3. **Complete the checklist** — Mark claims verified, facts current, accessibility checked, and legal wording reviewed before sign-off. **Why it matters:** Makes compliance review explicit and repeatable.
4. **Sign off** — Submit the reviewer sign-off. The Manager then makes the final approval, change-request, or rejection decision. **Why it matters:** Preserves separation of duties.

The Reviewer dashboard also reports turnaround, confirmed/dismissed outcomes, and overturned decisions.

Admin — maintain the workspace

Admins manage accounts, approved material, runtime configuration, and integrity tools. Admins cannot approve submissions.

Area	What to do	Why it matters
Users	Create or deactivate accounts; reset a password when needed.	Keeps access limited to active role accounts.
Approved Sources	Maintain guideline clauses, approved facts, terminology, and asset references.	Review quality depends on accurate source material.
Feedback and Insights	Inspect acceptance statistics and reset learned rewrite preferences if required.	Shows whether suggestions are being accepted or disputed.
Audit and Integrity	Review audit events and verify certificate/audit-chain consistency.	Helps detect changes to recorded workflow history.
System Status	Check database, source index, and LLM status; read the retrieval badge in the navigation.	Makes local setup and engine state visible.

The three Round 3 improvements

A — Guideline-Gap Insights

What it does: Measures dispute friction by clause and groups similar reasons. The manager/admin health page highlights risky wording and supports recompute and revision.
Why it matters: Important because repeated disagreements can point to unclear policy, not just individual mistakes.

B — Image Review and Format-Aware PDF Inspection

What it does: Uses OCR text boxes for images and page/line coordinates for PDFs. Reviewers can follow findings to the visual location; image checks compare against approved palette, logo, clear-space, and contrast rules when the required reference data exists.
Why it matters: Important because visual placement, claims, and readability can be missed by plain text review.

C — Reinforcement Feedback Loop

What it does: Learns from accepted fixes, edits, and disputes; ranks rewrites by predicted acceptance and suppresses a clause/span after three distinct disputes.
Why it matters: Important because human outcomes help reduce repeated, low-value suggestions.

Limits to remember: OCR and PDF coordinates can be approximate; approved asset metadata is needed for palette/logo checks. The app is a single local workspace, not a hosted multi-user service. See docs/LIMITATIONS.md for the full list.